

```
In [1]: import argparse
import os
import time
import shutil

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.backends.cudnn as cudnn

import torchvision
import torchvision.transforms as transforms
import sys

sys.path.insert(1, '../software')
from models import *

global best_prec
use_gpu = torch.cuda.is_available()
print('=> Building model...')

batch_size = 128
model_name = "final_P2_VGG16_quant"
model = final_P2_VGG16_quant()

print(model)

normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243, 0.247])

train_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        normalize,
    ]))
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size, shuffle=True, num_workers=2)

test_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
```

```
        normalize,
    ]))

testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size, shuffle=True)

print_freq = 100 # every 100 batches, accuracy printed. Here, each batch includes both training and validation data.
# CIFAR10 has 50,000 training data, and 10,000 validation data.

def train(trainloader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter()
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    model.train()

    end = time.time()
    for i, (input, target) in enumerate(trainloader):
        # measure data loading time
        data_time.update(time.time() - end)

        input, target = input.cuda(), target.cuda()

        # compute output
        output = model(input)
        loss = criterion(output, target)

        # measure accuracy and record loss
        prec = accuracy(output, target)[0]
        losses.update(loss.item(), input.size(0))
        top1.update(prec.item(), input.size(0))

        # compute gradient and do SGD step
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # measure elapsed time
        batch_time.update(time.time() - end)
        end = time.time()

        if i % print_freq == 0:
            print('Epoch: [{0}][{1}/{2}]\t'
                  'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
                  'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
                  'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
                  'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                      epoch, i, len(trainloader), batch_time=batch_time,
                      data_time=data_time, loss=losses, top1=top1))

def validate(val_loader, model, criterion):
    batch_time = AverageMeter()
```

```

losses = AverageMeter()
top1 = AverageMeter()

# switch to evaluate mode
model.eval()

end = time.time()
with torch.no_grad():
    for i, (input, target) in enumerate(val_loader):

        input, target = input.cuda(), target.cuda()

        # compute output
        output = model(input)
        loss = criterion(output, target)

        # measure accuracy and record loss
        prec = accuracy(output, target)[0]
        losses.update(loss.item(), input.size(0))
        top1.update(prec.item(), input.size(0))

        # measure elapsed time
        batch_time.update(time.time() - end)
        end = time.time()

        if i % print_freq == 0: # This Line shows how frequently print out the
            print('Test: [{0}/{1}]\t'
                  'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
                  'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
                  'Prec {top1.val:.3f}% ({top1.avg:.3f}%}'.format(
                      i, len(val_loader), batch_time=batch_time, loss=losses,
                      top1=top1))

    print(' * Prec {top1.avg:.3f}%'.format(top1=top1))
    return top1.avg

def accuracy(output, target, topk=(1,)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk)
    batch_size = target.size(0)

    _, pred = output.topk(maxk, 1, True, True)
    pred = pred.t()
    correct = pred.eq(target.view(1, -1).expand_as(pred))

    res = []
    for k in topk:
        correct_k = correct[:k].view(-1).float().sum(0)
        res.append(correct_k.mul_(100.0 / batch_size))
    return res

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):

```

```
self.reset()

def reset(self):
    self.val = 0
    self.avg = 0
    self.sum = 0
    self.count = 0

def update(self, val, n=1):
    self.val = val
    self.sum += val * n
    self.count += n
    self.avg = self.sum / self.count

def save_checkpoint(state, is_best, fdir):
    filepath = os.path.join(fdir, 'checkpoint.pth')
    torch.save(state, filepath)
    if is_best:
        shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))

def adjust_learning_rate(optimizer, epoch):
    """For resnet, the lr starts from 0.1, and is divided by 10 at 80 and 120 epoch
    adjust_list = [150, 225]
    #adjust_list = [100, 150, 225, 275]
    #adjust_list = [100, 150, 200, 250]
    #adjust_list = [100]
    if epoch in adjust_list:
        for param_group in optimizer.param_groups:
            param_group['lr'] = param_group['lr'] * 0.1

#model = nn.DataParallel(model).cuda()
#all_params = checkpoint['state_dict']
#model.load_state_dict(all_params, strict=False)
#criterion = nn.CrossEntropyLoss().cuda()
#validate(testloader, model, criterion)
```

```
=> Building model...
VGG_quant(
    (features): Sequential(
        (0): QuantConv2d(
            3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (2): ReLU(inplace=True)
        (3): QuantConv2d(
            64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (5): ReLU(inplace=True)
        (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (7): QuantConv2d(
            64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (9): ReLU(inplace=True)
        (10): QuantConv2d(
            128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (11): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (12): ReLU(inplace=True)
        (13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (14): QuantConv2d(
            128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (16): ReLU(inplace=True)
        (17): QuantConv2d(
            256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (18): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (19): ReLU(inplace=True)
        (20): QuantConv2d(
            256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (21): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
        (22): ReLU(inplace=True)
        (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
        (24): QuantConv2d(
```

```

        256, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
(25): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True, track_running_stats=
True)
    (26): ReLU(inplace=True)
    (27): QuantConv2d(
        16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
(28): ReLU(inplace=True)
(29): QuantConv2d(
        16, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
(30): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
=True)
    (31): ReLU(inplace=True)
    (32): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (33): QuantConv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
(34): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
=True)
    (35): ReLU(inplace=True)
    (36): QuantConv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
(37): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
=True)
    (38): ReLU(inplace=True)
    (39): QuantConv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
(40): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True, track_running_stats
=True)
    (41): ReLU(inplace=True)
    (42): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
    (43): AvgPool2d(kernel_size=1, stride=1, padding=0)
)
(classifier): Linear(in_features=512, out_features=10, bias=True)
)
Files already downloaded and verified
Files already downloaded and verified

```

In [2]: # This cell won't be given, but students will complete the training

```

lr = 1e-2
weight_decay = 1e-7
epochs = 300
best_prec = 0

#model = nn.DataParallel(model).cuda()

```

```

model.cuda()
criterion = nn.CrossEntropyLoss().cuda()
# SGD: Plateaus at about 86% on epoch 200, epoch 10 get 70%, epoch 40 get 80%, brea
# LR 1e-2, WD 1e-3, mom 0.9, epochs 120 (86.2% at 120, 1% increase in last ten)
# LR 4e-2, WD 4e-3, mom 0.9, epochs 120 (plateaus under 50% at around 40 epochs)
# LR 4e-2, WD 1e-3, mom 0.9, epochs 120 (77% by 30, 80% by 46, 82.7% by 120)
# LR 1e-2, WD 1e-4, mom 0.9, epochs 120 (small plateau at 50% for epoch 10, 58% by
# LR 5e-2, WD 1e-3, mom 0.9, epochs 120 (70% by 20, 80% by 50, 85% by 82, 86.5% by

#optimizer = torch.optim.SGD(model.parameters(), lr=Lr, momentum=0.9, weight_decay=
#scheduler = torch.optim.lr_scheduler.ReduceLROnPlateau(optimizer, threshold=1e-2,
# Does worse than SGD
#optimizer = torch.optim.RMSprop(model.parameters())
# Adam: (LR 1e-2, WD 1e-6, 70% by 10, 80% by 27, 84% by 100, 88% by 160, 88.4% by 2
# (LR 1e-2, WD 1e-7, 70% by 10, 80% by 20, 85% by 70, 86% by 128, 87 by 150, 88 by
# (LR 1e-2, WD 1e-8, 70% by 10, 80% by 23, 85% by 68, 86% by 104, 87 by 150, 88 by
optimizer = torch.optim.Adam(model.parameters(), lr=1e-2, weight_decay=1e-7)

# (LR 1e-3, WD 1e-7, 70% by 13, 80% by 21, 85% by 62, 86% by 86, 87% at 150, 88% at
# New Adjust, same as above: (70 at 12, 80 at 23, 85 by 53, 86 by 78, 87 by 91, 88
# Old Adjust, LR 5e-3, 1e-7
#optimizer = torch.optim.Adam(model.parameters(), lr=1e-2, weight_decay=1e-7)

#cudnn.benchmark = True

if not os.path.exists('result'):
    os.makedirs('result')
fdir = 'result/'+str(model_name)
if not os.path.exists(fdir):
    os.makedirs(fdir)

untrained = False

if (untrained):
    for epoch in range(0, epochs):
        #adjust_learning_rate(optimizer, epoch)

        train(trainloader, model, criterion, optimizer, epoch)

        # evaluate on test set
        print("Validation starts")
        prec = validate(testloader, model, criterion)

        # remember best precision and save checkpoint
        is_best = prec > best_prec
        best_prec = max(prec, best_prec)
        print('best acc: {:.1f}'.format(best_prec))
        save_checkpoint({
            'epoch': epoch + 1,
            'state_dict': model.state_dict(),
            'best_prec': best_prec,
            'optimizer': optimizer.state_dict(),
        }, is_best, fdir)

```

```
In [3]: # HW

# 1. Train with 4 bits for both weight and activation to achieve >90% accuracy
# 2. Find x_int and w_int for the 2nd convolution layer
# 3. Check the recovered psum has similar value to the un-quantized original psum
#     (such as example 1 in W3S2)
```

```
In [4]: PATH = "result/final_P2_VGG16_quant/model_best.pth.tar"
checkpoint = torch.load(PATH)
model.load_state_dict(checkpoint['state_dict'])
device = torch.device("cuda")

model.cuda()
model.eval()

test_loss = 0
correct = 0

with torch.no_grad():
    for data, target in testloader:
        data, target = data.to(device), target.to(device) # Loading to GPU
        output = model(data)
        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

test_loss /= len(testloader.dataset)

print('\nTest set: Accuracy: {} / {} ({:.0f}%)'.format(
    correct, len(testloader.dataset),
    100. * correct / len(testloader.dataset)))
```

Test set: Accuracy: 8705/10000 (87%)

```
In [5]: #send an input and grap the value by using prehook Like HW3
```

```
In [6]: class SaveOutput:
    def __init__(self):
        self.outputs = []
    def __call__(self, module, module_in):
        self.outputs.append(module_in)
    def clear(self):
        self.outputs = []

##### Save inputs from selected Layer #####
save_output = SaveOutput()

for layer in model.modules():
    if isinstance(layer, torch.nn.Conv2d):
        print("prehooked")
        layer.register_forward_pre_hook(save_output) ## Input for the module
#####

use_gpu = torch.cuda.is_available()
device = torch.device("cuda" if use_gpu else "cpu")
```

```

dataiter = iter(trainloader)
images, labels = next(dataiter)
images = images.to(device)
out = model(images)

print("1st convolution's input size:", save_output.outputs[0][0].size())
print("2nd convolution's input size:", save_output.outputs[1][0].size())

prehooked
1st convolution's input size: torch.Size([128, 3, 32, 32])
2nd convolution's input size: torch.Size([128, 64, 32, 32])

```

In [7]:

```

for index in range(len(save_output.outputs)):
    print(str(index) + ": " + str(save_output.outputs[index][0].size()))

```

```

0: torch.Size([128, 3, 32, 32])
1: torch.Size([128, 64, 32, 32])
2: torch.Size([128, 64, 16, 16])
3: torch.Size([128, 128, 16, 16])
4: torch.Size([128, 128, 8, 8])
5: torch.Size([128, 256, 8, 8])
6: torch.Size([128, 256, 8, 8])
7: torch.Size([128, 256, 4, 4])
8: torch.Size([128, 16, 4, 4])
9: torch.Size([128, 16, 4, 4])
10: torch.Size([128, 512, 2, 2])
11: torch.Size([128, 512, 2, 2])
12: torch.Size([128, 512, 2, 2])

```

In [8]:

```

w_bit = 4
weight_q = model.features[27].weight_q # quantized value is stored during the train
w_alpha = model.features[27].weight_quant.wgt_alpha # alpha is defined in your mod
w_delta = w_alpha / (2**((w_bit-1)-1)) # delta can be calculated by using alpha an
weight_int = weight_q / w_delta # w_int can be calculated by weight_q and w_delta
#print(weight_int) # you should see clean integer numbers

```

In [9]:

```

x_bit = 2
x = save_output.outputs[8][0] # input of the 2nd conv layer
x_alpha = model.features[27].act_alpha
x_delta = x_alpha / (2**((x_bit)-1))

act_quant_fn = act_quantization(x_bit) # define the quantization function
x_q = act_quant_fn(x, x_alpha) # create the quantized value for x

```

```
x_int = x_q / x_delta
#print(x_int) # you should see clean integer numbers
```

```
In [10]: conv_int = torch.nn.Conv2d(in_channels = 16, out_channels=16, kernel_size = 3, padding=1)
conv_int.weight = torch.nn.Parameter(Parameter(weight_int))

output_int = conv_int.forward(x_int)      # output_int can be calculated with conv_int
output_recovered = output_int * x_delta * w_delta  # recover with x_delta and w_delta

relu = torch.nn.ReLU()
relu_output_recovered = relu(output_recovered)
#print(output_recovered)
```

```
In [11]: next_input = save_output.outputs[9][0]

print(next_input.size())
print(relu_output_recovered.size())

torch.Size([128, 16, 4, 4])
torch.Size([128, 16, 4, 4])
```

```
In [12]: difference = abs( next_input - relu_output_recovered )
print(difference.mean()) ## It should be small, e.g., 2.3 in my trained model

tensor(1.8737e-05, device='cuda:0', grad_fn=<MeanBackward0>)
```

```
In [13]: print(x_int[0,:,:,:].size())

torch.Size([16, 4, 4])
```

```
In [78]: # act_int.size = torch.Size([128, 64, 32, 32])  <- batch_size, input_ch, ni, nj
a_int = x_int[0,:,:,:]
# a_int.size() = [64, 32, 32]

# conv_int.weight.size() = torch.Size([64, 64, 3, 3])  <- output_ch, input_ch, ki, kj
w_int = torch.reshape(weight_int, (weight_int.size(0), weight_int.size(1), -1))
# w_int.weight.size() = torch.Size([64, 64, 9])

padding = 1
stride = 1
row_num = 16 # row and column number
col_num = 8

nig = range(a_int.size(1)) ## ni group
njg = range(a_int.size(2)) ## nj group

icg = range(int(w_int.size(1))) ## input channel
ocg = range(int(w_int.size(0))) ## output channel

ic_tileg = range(int(len(icg)/row_num))
oc_tileg = range(int(len(ocg)/col_num))

kijg = range(w_int.size(2))
ki_dim = int(math.sqrt(w_int.size(2))) ## Kernel's 1 dim size

##### Padding before Convolution #####
a_pad = torch.zeros(len(icg), len(nig)+padding*2, len(nig)+padding*2).cuda()
```

```

# a_pad.size() = [64, 32+2pad, 32+2pad]
a_pad[ :, padding:padding+len(nijg), padding:padding+len(njg)] = a_int.cuda()
a_pad = torch.reshape(a_pad, (a_pad.size(0), -1))
# a_pad.size() = [64, (32+2pad)*(32+2pad)]


a_tile = torch.zeros(len(ic_tileg), row_num, a_pad.size(1)).cuda()
w_tile = torch.zeros(len(oc_tileg)*len(ic_tileg), col_num, row_num, len(kijg)).cuda

for ic_tile in ic_tileg:
    a_tile[ic_tile,:,:] = a_pad[ic_tile*row_num:(ic_tile+1)*row_num,:]

for ic_tile in ic_tileg:
    for oc_tile in oc_tileg:
        w_tile[ic_tile*len(oc_tileg) + oc_tile,:,:,:] = w_int[oc_tile*col_num:(oc_t

#####
p_nijg = range(a_pad.size(1)) ## psum nij group

psum = torch.zeros(len(ic_tileg), len(oc_tileg), col_num, len(p_nijg), len(kijg)).c

for kij in kijg:
    for ic_tile in ic_tileg:      # Tiling into array_sizeXarray_size array
        for oc_tile in oc_tileg:  # Tiling into array_sizeXarray_size array
            for nij in p_nijg:    # time domain, sequentially given input
                m = nn.Linear(col_num, row_num, bias=False)
                #m.weight = torch.nn.Parameter(w_int[oc_tile*array_size:(oc_til
                m.weight = torch.nn.Parameter(w_tile[len(oc_tileg)*ic_tile+oc_t
                psum[ic_tile, oc_tile, :, nij, kij] = m(a_tile[ic_tile,:,:nij]).c

```

```

In [79]: import math

a_pad_ni_dim = int(math.sqrt(a_pad.size(1))) # 32

o_hi_dim = int((a_pad_ni_dim - (ki_dim- 1) - 1)/stride + 1)
o_nijg = range(o_hi_dim**2)

print(len(o_nijg))

out = torch.zeros(len(ocg), len(o_nijg)).cuda()

### SFP accumulation ###
for o_nij in o_nijg:
    for kij in kijg:
        for ic_tile in ic_tileg:
            for oc_tile in oc_tileg:
                out[oc_tile*col_num:(oc_tile+1)*col_num, o_nij] = out[oc_tile*col_n
                psum[ic_tile, oc_tile, :, int(o_nij/o_hi_dim)*a_pad_ni_dim + o_nij%32
                ## 4th index = (int(o_nij/30)*32 + o_nij%30) + (int(kij/3)*32 + kij

```

```
In [80]: out_2D = torch.reshape(out, (out.size(0), o_ni_dim, -1)) # nij -> ni & nj
difference = (out_2D - output_int[0,:,:,:])
print(difference.abs().sum())
tensor(7.1526e-07, device='cuda:0', grad_fn=<SumBackward0>)
```

```
In [105...]: if not os.path.exists("P2_Files"):
    os.makedirs("P2_Files")
if not os.path.exists("P2_Files/Tile0"):
    os.makedirs("P2_Files/Tile0")
if not os.path.exists("P2_Files/Tile1"):
    os.makedirs("P2_Files/Tile1")
```

```
In [106...]: ### show this cell partially. The following cells should be printed by students ####
tile_id = 0
nij = 200 # just a random number
X = a_tile[tile_id,:,:] # [tile_num, array row num, time_steps]

bit_precision = 2
file = open('P2_Files/activation.txt', 'w') #write to file
file.write('#time0row7[msb-lsb],time0row6[msb-lst],...,time0row0[msb-lst]\#\n')
file.write('#time1row7[msb-lsb],time1row6[msb-lst],...,time1row0[msb-lst]\#\n')
file.write('#.....#\n')

for i in range(X.size(1)): # time step
    for j in range(X.size(0)): # row #
        X_bin = '{0:02b}'.format(round(X[X.size(0)-1-j,i].item()))
        for k in range(bit_precision):
            file.write(X_bin[k])
        #file.write(' ') # for visibility with blank between words, you can use
        file.write('\n')
file.close() #close file
```

```
In [111...]: ### Complete this cell ####
kij = 0
array_row_num = 8

bit_precision = 4
for tile_id in range(2):
    for kij in range(9):
        W = w_tile[tile_id,:,:,:,kij] # w_tile[tile_num, array col num, array row num]
        file = open('P2_Files/Tile' + str(tile_id) + '/weight_' + str(kij) + '.txt'
        file.write('#col0row14[msb-lsb],col0row12[msb-lst],...,col0row0[msb-lst]\#\n')
        file.write('#col0row15[msb-lsb],col0row13[msb-lst],...,col0row1[msb-lst]\#\n'
#        file.write('#col1row7[msb-lsb],col1row6[msb-lst],...,col1row0[msb-lst]\#\n'
        file.write('#.....#\n')

        for i in range(W.size(0)): # col
            for m in range(2):
                for j in range(array_row_num): # row #

                    #print(W.size(1)-2*(j+1)+m)
                    W_bin = '{0:04b}'.format(round(W[i,W.size(1)-2*(j+1)+m].item()))
                    for k in range(bit_precision):
                        file.write(W_bin[k])
```

```

        #file.write(' ') # for visibility with blank between words, you can
        file.write('\n')
        file.close() #close file
print(W.size())

torch.Size([8, 16])

```

```

In [112... ]### Complete this cell ###
ic_tile_id = 0

kij = 0
nij = 200
# psum[Len(ic_tileg), Len(oc_tileg), array_size, len(p_nijg), Len(kijg)]


bit_precision = 16

for oc_tile_id in range(2):
    for kij in range(9):
        psum_tile = psum[ic_tile_id,oc_tile_id,:,:,:,kij]
        file = open('P2_Files/Tile' + str(oc_tile_id) + '/psum_' + str(kij) + '.txt'
        file.write('#time0col7[msb-lsb],time0col6[msb-lst],...,time0col0[msb-lst]#\n')
        file.write('#time1col7[msb-lsb],time1col6[msb-lst],...,time1col0[msb-lst]#\n')
        file.write('#.....#\n')

        for i in range(psum_tile.size(1)): # time step
            for j in range(psum_tile.size(0)): # col #
                #psum_bin = '{0:016b}'.format(round(psum_tile[psum_tile.size(0)-1-j,i].item()))
                curr_psum = round(psum_tile[psum_tile.size(0)-1-j,i].item())
                #if (i == 7 and kij == 0):
                #    print(curr_psum)
                if (curr_psum < 0):
                    #if (i == 7 and kij == 0):
                    #    print('{0:016b}'.format(curr_psum))
                    curr_psum = curr_psum + 2**bit_precision
                    #if (i == 7 and kij == 0):
                    #    print('{0:016b}'.format(curr_psum))
                psum_bin = '{0:016b}'.format(curr_psum)

                for k in range(bit_precision):
                    file.write(psum_bin[k])
                    #file.write(' ') # for visibility with blank between words, you can
                    file.write('\n')
        file.close() #close file

```

```

In [119... ]### Complete this cell ###

bit_precision = 16
# out is array of size columns x len(o_nijg)

for out_tile in range(2):
    file = open('P2_Files/Tile' + str(out_tile) + '/out.txt', 'w') #write to file
    file.write('#time0col7[msb-lsb],time0col6[msb-lst],...,time0col0[msb-lst]#\n')
    file.write('#time1col7[msb-lsb],time1col6[msb-lst],...,time1col0[msb-lst]#\n')
    file.write('#.....#\n')

```

```

for i in range(out.size(1)): # time step
    for j in range(col_num): # col #
        out_bin = '{0:016b}'.format(0 if (round(out[col_num*(out_tile+1)-1-j,i])
        for k in range(bit_precision):
            file.write(out_bin[k])
        #file.write(' ') # for visibility with blank between words, you can use
        file.write('\n')
file.close() #close file

```

In [114... print(X[:,7])

```
tensor([0., 0., 0., 0., 2., 0., 0., 1., 1., 1., 0., 0., 0., 1., 1., 1., 0.],
      device='cuda:0', grad_fn=<SelectBackward0>)
```

In [115... print(x_int.size())

```
torch.Size([128, 16, 4, 4])
```

In [116... print(w_int.size())

```
torch.Size([16, 16, 9])
```

In [117... print(w_int[:, :, 0])

```
tensor([[ 2.,  1., -3., -4.,  0., -0., -0., -4.,  1.,  2.,  1., -2., -1.,  1.,
         -2., -1.],
       [-0., -1.,  4., -2., -2., -2.,  1.,  3.,  0., -5., -3.,  4.,  1., -4.,
        2.,  0.],
       [ 1., -3., -1., -0., -2.,  3.,  0., -1., -2., -1.,  3.,  0., -0.,  1.,
        1., -0.],
       [ 4.,  4.,  1.,  4., -1.,  2.,  2., -1., -1., -1., -1., -0.,  2.,  0.,
        1.,  1.],
       [ 0., -5.,  1.,  1., -2.,  1., -0.,  1., -5.,  1., -3., -2.,  3.,  2.,
        1.,  4.],
       [ 1.,  1.,  1., -1., -2., -1., -0.,  3.,  1., -3.,  2., -1., -1., -1.,
        0.,  0.],
       [ 3.,  1.,  0., -2., -1., -2.,  6., -1.,  0., -0., -2., -1.,  1.,  1.,
        -3.,  1.],
       [ 1.,  1., -4.,  2.,  2.,  0., -2.,  1., -4.,  2.,  1.,  0.,  2., -2.,
        3.,  1.],
       [-0.,  1., -2.,  1., -2.,  3., -2., -0., -5., -2., -2.,  1.,  1.,  1.,
        1.,  1.],
       [ 1., -1., -1., -2.,  1., -0., -4., -2.,  3., -2.,  0.,  1., -1., -2.,
        -1., -5.],
       [-1.,  1., -0., -2., -1.,  2., -3.,  0.,  1.,  1.,  2.,  0., -4., -1.,
        2., -1.],
       [ 1.,  1.,  2., -0.,  1., -0.,  0.,  2.,  0.,  2., -2., -4.,  0.,  3.,
        -3., -1.],
       [-2., -3.,  2.,  2.,  1.,  2., -2.,  1.,  1., -0.,  1., -2., -1., -1.,
        4., -1.],
       [-4., -0.,  2., -4., -1., -0.,  1., -3., -1., -3.,  2., -1., -5., -1.,
        0.,  0.],
       [-3.,  0.,  0.,  0., -4., -2., -0.,  1., -3.,  0.,  0., -1., -0., -2.,
        -1., -1.],
       [-0., -0., -3.,  1.,  2., -0.,  1.,  2.,  0., -0., -6., -2.,  2., -2.,
        3., -3.]], device='cuda:0', grad_fn=<SelectBackward0>)
```

```
In [93]: print(w_tile[0,:,:,:0])
```

```
tensor([[ 2.,  1., -3., -4.,  0., -0., -0., -4.,  1.,  2.,  1., -2., -1.,  1.,
         -2., -1.],
       [-0., -1.,  4., -2., -2., -2.,  1.,  3.,  0., -5., -3.,  4.,  1., -4.,
        2.,  0.],
       [ 1., -3., -1., -0., -2.,  3.,  0., -1., -2., -1.,  3.,  0., -0.,  1.,
        1., -0.],
       [ 4.,  4.,  1.,  4., -1.,  2.,  2., -1., -1., -1., -1., -0.,  2.,  0.,
        1.,  1.],
       [ 0., -5.,  1.,  1., -2.,  1., -0.,  1., -5.,  1., -3., -2.,  3.,  2.,
        1.,  4.],
       [ 1.,  1.,  1., -1., -2., -1., -0.,  3.,  1., -3.,  2., -1., -1., -1.,
        0.,  0.],
       [ 3.,  1.,  0., -2., -1., -2.,  6., -1.,  0., -0., -2., -1.,  1.,  1.,
        -3.,  1.],
       [ 1.,  1., -4.,  2.,  2.,  0., -2.,  1., -4.,  2.,  1.,  0.,  2., -2.,
        3.,  1.]], device='cuda:0', grad_fn=<SelectBackward0>)
```

```
In [94]: print(w_tile.size())
```

```
torch.Size([2, 8, 16, 9])
```

```
In [ ]:
```